

OpenCity AI Whitepaper

OpenCity AI: A Lightweight, Open-Source Municipal Knowledge Retrieval Platform

Abstract

OpenCity AI is an open-source, multi-tenant platform that helps municipalities and civic technology communities improve access to public information using Retrieval-Augmented Generation (RAG). It is designed to complement existing city systems by providing a document-grounded retrieval and answer layer with source citations and predictable operating cost.

1. Problem

Cities publish service procedures, policy pages, and regulations across many systems. Citizens and staff often face information-access friction because the data is distributed, unevenly indexed, and updated over time.

2. Proposed System

```
``mermaid
flowchart LR
A[Citizen or Staff Query] --> B[FastAPI Gateway]
B --> C[Retriever]
C --> D[Qdrant city_id filter]
D --> E[Context Chunks]
E --> F[Open-Source LLM]
F --> G[Answer with Citations]
``
```

3. Multi-Tenant Design

```
``mermaid
flowchart TD
A[City A] --> S[Shared API]
B[City B] --> S
C[City C] --> S
S --> V[Shared Qdrant]
V --> F1[city_id=city_a]
V --> F2[city_id=city_b]
V --> F3[city_id=city_c]
``
```

Each query is filtered by 'city_id', so retrieval stays scoped to the correct city corpus.

4. Knowledge Updates

1. Fetch document sources.
2. Compute content hash.
3. If changed: parse, chunk, embed, and upsert vectors.

4. Queries immediately use latest indexed version.

5. Cost Model (Pilot)

Component	Typical Spec	Monthly Cost (USD)
--- --- ---		
App + retrieval server	4-8 vCPU, 8-16GB RAM	40-100
Storage volumes	Postgres + Qdrant	5-20
LLM inference	Open-source local runtime	Included in infra
Total (pilot range)	Single-node deployment	45-120

6. Limitations

- Assistive retrieval system, not autonomous city decision-making.
- Accuracy depends on source quality and update discipline.
- Human verification is still required for high-stakes actions.

7. Conclusion

OpenCity AI provides practical civic value when positioned as knowledge infrastructure: transparent retrieval, source-cited responses, and low-cost deployment that can integrate with existing municipal systems.